

```
pip install numpy pandas matplotlib seaborn scikit-learn tensorflow torch torchvision torchaudio opencv-python pillow alumentat
```

```
Requirement already satisfied: numpy in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (1.2
Requirement already satisfied: pandas in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (2.
Requirement already satisfied: matplotlib in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
Requirement already satisfied: seaborn in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (0
Requirement already satisfied: scikit-learn in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packag
Requirement already satisfied: tensorflow in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
Requirement already satisfied: torch in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (2.5
Requirement already satisfied: torchvision in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-package
Requirement already satisfied: torchaudio in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
Requirement already satisfied: opencv-python in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packa
Requirement already satisfied: pillow in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (11
Requirement already satisfied: alumentations in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pack
Requirement already satisfied: grad-cam in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (
Requirement already satisfied: python-dateutil>=2.8.2 in /Users/jishnunarasisimhamoorthy/Library/Python/3.12/lib/python/site-p
Requirement already satisfied: pytz>=2020.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packag
Requirement already satisfied: tzdata>=2022.7 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pack
Requirement already satisfied: contourpy>=1.0.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pa
Requirement already satisfied: cycycler>=0.10 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packag
Requirement already satisfied: fonttools>=4.22.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-p
Requirement already satisfied: kiwisolver>=1.3.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-p
Requirement already satisfied: packaging>=20.0 in /Users/jishnunarasisimhamoorthy/Library/Python/3.12/lib/python/site-packages
Requirement already satisfied: pyparsing>=2.3.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pa
Requirement already satisfied: scipy>=1.6.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packag
Requirement already satisfied: joblib>=1.2.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packa
Requirement already satisfied: threadpoolctl>=3.1.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/sit
Requirement already satisfied: absl-py>=1.0.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pack
Requirement already satisfied: astunparse>=1.6.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-p
Requirement already satisfied: flatbuffers>=24.3.25 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/sit
Requirement already satisfied: gast!=0.5.0,!=0.5.1,!=0.5.2,>=0.2.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib
Requirement already satisfied: google-pasta>=0.1.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site
Requirement already satisfied: libclang>=13.0.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pa
Requirement already satisfied: opt-einsum>=2.3.2 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-p
Requirement already satisfied: protobuf!=4.21.0,!=4.21.1,!=4.21.2,!=4.21.3,!=4.21.4,!=4.21.5,<6.0.0dev,>=3.20.3 in /Library/
Requirement already satisfied: requests<3,>=2.21.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site
Requirement already satisfied: setuptools in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages
Requirement already satisfied: six>=1.12.0 in /Users/jishnunarasisimhamoorthy/Library/Python/3.12/lib/python/site-packages (fr
Requirement already satisfied: termcolor>=1.1.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pa
Requirement already satisfied: typing-extensions>=3.6.6 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12
Requirement already satisfied: wrapt>=1.11.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packa
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site
Requirement already satisfied: tensorboard<2.19,>=2.18 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
Requirement already satisfied: keras>=3.5.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packag
Requirement already satisfied: h5py>=3.11.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packag
Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
Requirement already satisfied: filelock in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (
Requirement already satisfied: networkx in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (
Requirement already satisfied: Jinja2 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (fr
Requirement already satisfied: fsspec in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (fr
Requirement already satisfied: sympy==1.13.1 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packa
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-
Requirement already satisfied: PyYAML in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (fr
Requirement already satisfied: pydantic>=2.9.2 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pac
Requirement already satisfied: albucore==0.0.21 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pa
Requirement already satisfied: eval-type-backport in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-
Requirement already satisfied: opencv-python-headless>=4.9.0.80 in /Library/Frameworks/Python.framework/Versions/3.12/lib/py
Requirement already satisfied: stringzilla>=3.10.4 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site
Requirement already satisfied: simsimd>=5.9.2 in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-pack
Requirement already satisfied: ttach in /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages (fro
```

Core Libraries

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Machine Learning Libraries

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, roc_curve, confusion_matrix, classification_report
```

Deep Learning Libraries

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

Visualization

```
import cv2
from tensorflow.keras.utils import plot_model
import albumentations as A
```

```
# Grad-CAM
from tensorflow.keras.models import Model
from tensorflow.keras.preprocessing.image import img_to_array, load_img
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.applications.resnet import preprocess_input
```

```
↗ /Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/albumentations/check_version.py:51: UserWarn
data = fetch_version_info()
```

```
def check_gpu():
    print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
    if tf.test.gpu_device_name():
        print("GPU is enabled:", tf.test.gpu_device_name())
    else:
        print("GPU is not enabled.")
```

```
check_gpu()
```

```
↗ Num GPUs Available: 0
GPU is not enabled.
```

```
# Define paths to the dataset
base_path = "archive/sample"
labels_path = os.path.join(base_path, "sample_labels.csv")
images_path = os.path.join(base_path, "images")
```

Data Cleaning and Loading

```
# Load the CSV file
data = pd.read_csv(labels_path)

# Add image paths to the dataframe
data['image_path'] = data['Image Index'].apply(lambda x: os.path.join(images_path, x))

# Clean Patient Age
# Extract numeric values and convert to integer
data['Patient Age'] = data['Patient Age'].str.extract('(\d+)').astype(float).astype(int)

# Replace missing labels with an empty string
data['Finding Labels'] = data['Finding Labels'].fillna("")

# List of disease classes
disease_classes = [
    'Atelectasis', 'Cardiomegaly', 'Effusion', 'Infiltration',
    'Mass', 'Nodule', 'Pneumonia', 'Pneumothorax',
    'Consolidation', 'Edema', 'Emphysema', 'Fibrosis',
    'Pleural_Thickening', 'Hernia'
]

# Generate binary columns for each disease class
for disease in disease_classes:
    data[disease] = data['Finding Labels'].apply(lambda x: 1 if disease in x else 0)

# Check the processed data
print(data.head())
```

```
↗
```

	Image Index	Finding Labels \
0	00000013_005.png	Emphysema Infiltration Pleural_Thickening Pneu...
1	00000013_026.png	Cardiomegaly Emphysema
2	00000017_001.png	No Finding
3	00000030_001.png	Atelectasis
4	00000032_001.png	Cardiomegaly Edema Effusion

	Follow-up #	Patient ID	Patient Age	Patient Gender	View Position \
0	5	13	60	M	AP
1	26	13	57	M	AP
2	1	17	77	M	AP
3	1	30	79	M	PA
4	1	32	55	F	AP

	OriginalImageWidth	OriginalImageHeight	OriginalImagePixelSpacing_x	...	\
0	3056	2544	0.139	...	
1	2500	2048	0.168	...	
2	2500	2048	0.168	...	
3	2992	2991	0.143	...	
4	2500	2048	0.168	...	

	Mass	Nodule	Pneumonia	Pneumothorax	Consolidation	Edema	Emphysema	\
0	0	0	0	1	0	0	1	
1	0	0	0	0	0	0	1	
2	0	0	0	0	0	0	0	
3	0	0	0	0	0	0	0	
4	0	0	0	0	0	1	0	

	Fibrosis	Pleural_Thickening	Hernia
0	0	1	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

```
[5 rows x 26 columns]
```

```
<=:9: SyntaxWarning: invalid escape sequence '\d'
```

```
<=:9: SyntaxWarning: invalid escape sequence '\d'
```

```
/var/folders/wn/n2fx_jqn07z9z11jmc33c0000000gn/T/ipykernel_62307/505373478.py:9: SyntaxWarning: invalid escape sequence '\d'
data['Patient Age'] = data['Patient Age'].str.extract('(\d+)').astype(float).astype(int)
```

Checking for Null Values

```
print(data.isnull().sum())
```

```

Image Index      0
Finding Labels   0
Follow-up #      0
Patient ID       0
Patient Age      0
Patient Gender   0
View Position    0
OriginalImageWidth  0
OriginalImageHeight 0
OriginalImagePixelSpacing_x 0
OriginalImagePixelSpacing_y 0
image_path       0
Atelectasis      0
Cardiomegaly     0
Effusion         0
Infiltration     0
Mass             0
Nodule           0
Pneumonia        0
Pneumothorax     0
Consolidation    0
Edema            0
Emphysema        0
Fibrosis         0
Pleural_Thickening 0
Hernia           0
dtype: int64

```

Checking for Label Consistency

```
unique_labels = set('|'.join(data['Finding Labels']).split('|'))
print("Unique labels in dataset:", unique_labels)
```

```
Unique labels in dataset: {'Pleural_Thickening', 'Emphysema', 'Pneumothorax', 'Pneumonia', 'Infiltration', 'Mass', 'Atelecta
```

Making sure columns only contain 1 and 0

```
print(data[disease_classes].sum())
print(data[disease_classes].head())
```

```

Atelectasis      508
Cardiomegaly     141
Effusion         644
Infiltration     967

```

```

Mass                284
Nodule              313
Pneumonia           62
Pneumothorax        271
Consolidation       226
Edema               118
Emphysema           127
Fibrosis            84
Pleural_Thickening  176
Hernia              13
dtype: int64
Atelectasis  Cardiomegaly  Effusion  Infiltration  Mass  Nodule  Pneumonia  \
0            0            0          0            1    0      0          0
1            0            1          0            0    0    0      0
2            0            0          0            0    0    0      0
3            1            0          0            0    0    0      0
4            0            1          1            0    0    0      0

Pneumothorax  Consolidation  Edema  Emphysema  Fibrosis  \
0            1            0      0          1          0
1            0            0      0          1          0
2            0            0      0          0          0
3            0            0      0          0          0
4            0            0      1          0          0

Pleural_Thickening  Hernia
0                  1      0
1                  0      0
2                  0      0
3                  0      0
4                  0      0

```

Check for Multi-Label Rows : Making sure there are rows with multiple classifications.

```

multi_label_rows = data[disease_classes].sum(axis=1)
print(f"Number of multi-label rows: {sum(multi_label_rows > 1)}")

```

➦ Number of multi-label rows: 980

```
print(f"Dataset size after cleaning: {len(data)}")
```

➦ Dataset size after cleaning: 5606

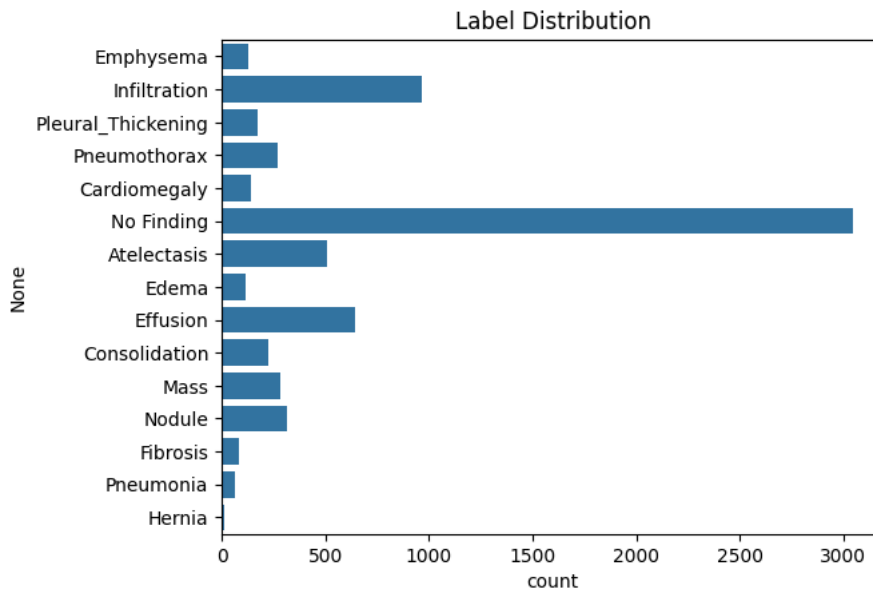
EDA

```

def visualize_label_distribution(data):
    """
    Visualize the distribution of labels.
    """
    all_labels = [label for sublist in data['Finding Labels'].str.split('|') for label in sublist]
    sns.countplot(y=pd.Series(all_labels))
    plt.title("Label Distribution")
    plt.show()

```

```
visualize_label_distribution(data)
```



```
def display_one_image_per_class(data):
    """
    Display one random image for each of the 14 unique disease classes in the dataset.
    """
    # List of all 14 possible disease classes
    all_classes = [
        'Atelectasis', 'Cardiomegaly', 'Effusion', 'Infiltration',
        'Mass', 'Nodule', 'Pneumonia', 'Pneumothorax',
        'Consolidation', 'Edema', 'Emphysema', 'Fibrosis',
        'Pleural_Thickening', 'Hernia'
    ]

    for disease in all_classes:
        # Filter data to get rows that include the current disease
        class_data = data[data['Finding Labels'].str.contains(disease, na=False)]


        if not class_data.empty:
            # Select one random image from this class
            sample = class_data.sample(1).iloc[0]

            # Load and display the image
            img = load_img(sample['image_path'], target_size=(224, 224))
            plt.imshow(img)
            plt.title(f"Class: {disease}")
            plt.axis('off')
            plt.show()
        else:
            print(f"No images found for class: {disease}")

    # Call the function
    display_one_image_per_class(data)
```



Class: Atelectasis



Class: Cardiomegaly



Class: Effusion




Class: Infiltration






Class: Mass



Class: Nodule



Class: Pneumonia



Class: Pneumothorax



Class: Consolidation



Class: Edema



Class: Emphysema





Class: Fibrosis



Class: Pleural_Thickening



Class: Hernia



Train-Test Split

```
from sklearn.model_selection import train_test_split

# Split the dataset into train and test sets (80% train, 20% test)
train_data, test_data = train_test_split(data, test_size=0.2, random_state=42)

print(f"Training set size: {len(train_data)}")
print(f"Testing set size: {len(test_data)}")
```

↗ Training set size: 4484
Testing set size: 1122

Pre-Processing for Neural Network

```
datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=10,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    zoom_range=0.1,
    horizontal_flip=True,
    fill_mode='nearest'
)

# Training generator
train_gen = datagen.flow_from_dataframe(
    train_data,
    x_col='image_path',
    y_col=disease_classes,
    target_size=(224, 224),
    color_mode='grayscale', # Load images as grayscale
    class_mode='raw', # Multi-label classification
    batch_size=32
)

# Testing generator
test_gen = datagen.flow_from_dataframe(
    test_data,
    x_col='image_path',
    y_col=disease_classes,
    target_size=(224, 224),
    color_mode='grayscale', # Load images as grayscale
    class_mode='raw', # Multi-label classification
    batch_size=32
)
```

↗ Found 4484 validated image filenames.
Found 1122 validated image filenames.

Base CNN Model

```
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.models import Model

# Define the Functional API Model
input_layer = Input(shape=(224, 224, 1), name='input_layer')
x = Conv2D(32, (3, 3), activation='relu', name='conv2d_1')(input_layer)
x = MaxPooling2D((2, 2), name='max_pooling2d_1')(x)
x = Conv2D(64, (3, 3), activation='relu', name='conv2d_2')(x)
x = MaxPooling2D((2, 2), name='max_pooling2d_2')(x)
x = Conv2D(128, (3, 3), activation='relu', name='conv2d_3')(x)
x = MaxPooling2D((2, 2), name='max_pooling2d_3')(x)
x = Flatten(name='flatten')(x)
x = Dense(128, activation='relu', name='dense_1')(x)
x = Dropout(0.5, name='dropout')(x)
output_layer = Dense(len(disease_classes), activation='sigmoid', name='output_layer')(x)

CNN_model = Model(inputs=input_layer, outputs=output_layer, name='CNN_Model')
```

```
# Compile the model
CNN_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Summary of the model
CNN_model.summary()
```

Model: "CNN_Model"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 224, 224, 1)	0
conv2d_1 (Conv2D)	(None, 222, 222, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_2 (Conv2D)	(None, 109, 109, 64)	18,496
max_pooling2d_2 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_3 (Conv2D)	(None, 52, 52, 128)	73,856
max_pooling2d_3 (MaxPooling2D)	(None, 26, 26, 128)	0
flatten (Flatten)	(None, 86528)	0
dense_1 (Dense)	(None, 128)	11,075,712
dropout (Dropout)	(None, 128)	0
output_layer (Dense)	(None, 14)	1,806

Total params: 11,170,190 (42.61 MB)
 Trainable params: 11,170,190 (42.61 MB)
 Non-trainable params: 0 (0.00 B)

```
from tensorflow.keras.callbacks import EarlyStopping
```

```
# Define the EarlyStopping callback
early_stopping = EarlyStopping(
    monitor='val_loss', # Metric to monitor
    patience=3,         # Number of epochs with no improvement
    restore_best_weights=True # Restore model weights from the best epoch
)
```

```
# Train the model
history = CNN_model.fit(
    train_gen,          # Training data generator
    validation_data=test_gen, # Validation data generator
    epochs=5,          # Adjust the number of epochs
    callbacks=[early_stopping], # Early stopping
    verbose=1
)
```

```
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/keras/src/trainers/data_adapters/py_dataset_
self._warn_if_super_not_called()
Epoch 1/5
141/141 ————— 107s 754ms/step - accuracy: 0.1662 - loss: 0.2607 - val_accuracy: 0.1176 - val_loss: 0.1819
Epoch 2/5
141/141 ————— 101s 717ms/step - accuracy: 0.1710 - loss: 0.1972 - val_accuracy: 0.1176 - val_loss: 0.1797
Epoch 3/5
141/141 ————— 105s 740ms/step - accuracy: 0.1833 - loss: 0.1965 - val_accuracy: 0.1176 - val_loss: 0.1796
Epoch 4/5
141/141 ————— 108s 762ms/step - accuracy: 0.1630 - loss: 0.1891 - val_accuracy: 0.1176 - val_loss: 0.1786
Epoch 5/5
141/141 ————— 106s 758ms/step - accuracy: 0.1781 - loss: 0.1858 - val_accuracy: 0.1194 - val_loss: 0.1790
```

```
def generate_gradcam_heatmap(model, last_conv_layer_name, img_array, class_index=None):
    grad_model = tf.keras.models.Model(
        [model.input],
        [model.get_layer(last_conv_layer_name).output, model.output]
    )
```

```
    with tf.GradientTape() as tape:
        conv_outputs, predictions = grad_model(img_array)
        if class_index is None:
            class_index = tf.argmax(predictions[0])
        class_channel = predictions[:, class_index]
```

```

grads = tape.gradient(class_channel, conv_outputs)
if grads is None:
    raise ValueError("Gradients could not be computed. Check model compatibility.")

# Debug gradients
print("Gradients shape:", grads.shape)
print("Gradients min:", tf.reduce_min(grads).numpy())
print("Gradients max:", tf.reduce_max(grads).numpy())

pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))
conv_outputs = conv_outputs[0]
heatmap = tf.reduce_sum(tf.multiply(pooled_grads, conv_outputs), axis=-1)

# Normalize the heatmap
heatmap = np.maximum(heatmap, 0)
heatmap = heatmap / np.max(heatmap) if np.max(heatmap) != 0 else np.zeros_like(heatmap)
return heatmap

# heatmap = np.maximum(heatmap, 0)
# max_value = np.max(heatmap)
# heatmap = heatmap / max_value if max_value > 1e-10 else np.zeros_like(heatmap)

import os
import random
from tensorflow.keras.preprocessing.image import load_img, img_to_array
import numpy as np

# Specify the image directory
image_dir = "archive/sample/images" # Replace with your actual directory

# List all image files in the directory
image_files = [os.path.join(image_dir, f) for f in os.listdir(image_dir) if f.endswith((".png", ".jpg", ".jpeg"))]

# Randomly select an image
img_path = random.choice(image_files)
print(f"Selected Image: {img_path}")

➡ Selected Image: archive/sample/images/00001301_021.png

# Load the selected image
img = load_img(img_path, target_size=(224, 224), color_mode="grayscale")

# Convert the image to an array and normalize it
img_array = np.expand_dims(img_to_array(img) / 255.0, axis=0) # Normalize to [0, 1]
print("Image shape:", img_array.shape) # Should print (1, 224, 224, 1) for grayscale

➡ Image shape: (1, 224, 224, 1)

# Predict the class for the selected image using the model
predictions = CNN_model.predict(img_array)
predicted_class_index = np.argmax(predictions[0]) # Get the index of the predicted class
predicted_class = disease_classes[predicted_class_index] # Use the index to find the class name

print(f"Predicted Class: {predicted_class}")

➡ 1/1 ————— 0s 27ms/step
Predicted Class: Infiltration

for idx, layer in enumerate(CNN_model.layers):
    if hasattr(layer, 'output_shape'):
        output_shape = layer.output_shape
    elif hasattr(layer, 'batch_input_shape'): # Handle InputLayer specifically
        output_shape = layer.batch_input_shape
    else:
        output_shape = "Unknown"
    print(f"Layer {idx}: {layer.name}, Output shape: {output_shape}")

➡ Layer 0: input_layer, Output shape: Unknown
Layer 1: conv2d_1, Output shape: Unknown
Layer 2: max_pooling2d_1, Output shape: Unknown
Layer 3: conv2d_2, Output shape: Unknown
Layer 4: max_pooling2d_2, Output shape: Unknown
Layer 5: conv2d_3, Output shape: Unknown
Layer 6: max_pooling2d_3, Output shape: Unknown

```

```

Layer 7: flatten, Output shape: Unknown
Layer 8: dense_1, Output shape: Unknown
Layer 9: dropout, Output shape: Unknown
Layer 10: output_layer, Output shape: Unknown

```

```
# Generate and display the heatmap
```

```
last_conv_layer_name = "conv2d_3" # Update this based on your model
heatmap = generate_gradcam_heatmap(CNN_model, last_conv_layer_name, img_array)
```

```
↗ Gradients shape: (1, 52, 52, 128)
```

```
Gradients min: -0.022241646
```

```
Gradients max: 0.029697776
```

```
# Resize the heatmap to match the image size
```

```
heatmap_resized = cv2.resize(heatmap, (224, 224))
```

```
# Overlay the heatmap on the original image
```

```
superimposed_img = cv2.addWeighted(img_to_array(img) / 255.0, 0.6, heatmap_resized, 0.4, 0)
```

```
# Get prediction probabilities for each class
```

```
predicted_probabilities = predictions[0]
```

```
# Print predicted probabilities for all classes
```

```
print("Prediction Probabilities:")
```

```
for i, class_name in enumerate(disease_classes):
```

```
    print(f"{class_name}: {predicted_probabilities[i]:.2f}")
```

```
# Get the predicted class label
```

```
predicted_class_index = np.argmax(predicted_probabilities) # Index of the predicted class
```

```
predicted_class = disease_classes[predicted_class_index] # Class label from disease_classes list
```

```
predicted_probability = predicted_probabilities[predicted_class_index] # Probability of the predicted class
```

```
# Plot the original image and Grad-CAM heatmap side by side
```

```
plt.figure(figsize=(12, 6))
```

```
# Original Image
```

```
plt.subplot(1, 2, 1)
```

```
plt.imshow(img_to_array(img) / 255.0, cmap='gray') # Display the original image in grayscale
```

```
plt.axis('off')
```

```
plt.title("Original Image")
```

```
# Grad-CAM Heatmap Overlayed Image
```

```
plt.subplot(1, 2, 2)
```

```
plt.imshow(superimposed_img, cmap='jet') # Display the heatmap overlay
```

```
plt.axis('off')
```

```
plt.title(f'Grad-CAM Heatmap\nPredicted Class: {predicted_class} ({predicted_probability:.2f})')
```



```
plt.show()
```

```

Prediction Probabilities:
Atelectasis: 0.15
Cardiomegaly: 0.05
Effusion: 0.19
Infiltration: 0.28
Mass: 0.09
Nodule: 0.08
Pneumonia: 0.02
Pneumothorax: 0.08
Consolidation: 0.07
Edema: 0.04
Emphysema: 0.05
Fibrosis: 0.03
Pleural_Thickening: 0.06
Hernia: 0.01

```

Original Image

Grad-CAM Heatmap
Predicted Class: Infiltration (0.28)

```

## Resize the heatmap to match the image size
# heatmap_resized = cv2.resize(heatmap, (224, 224))

## Overlay the heatmap on the original image
# superimposed_img = cv2.addWeighted(img_to_array(img) / 255.0, 0.6, heatmap_resized, 0.4, 0)

## Get prediction probabilities for each class
# predicted_probabilities = predictions[0]

## Print predicted probabilities for all classes
# print("Prediction Probabilities:")
# for i, class_name in enumerate(disease_classes):
#     print(f"{class_name}: {predicted_probabilities[i]:.2f}")

## Get the predicted class label
# predicted_class_index = np.argmax(predicted_probabilities) # Index of the predicted class
# predicted_class = disease_classes[predicted_class_index] # Class label from disease_classes list
# predicted_probability = predicted_probabilities[predicted_class_index] # Probability of the predicted class

## Display the image with heatmap, predicted class, and its probability
# plt.imshow(superimposed_img, cmap='jet')
# plt.axis('off')
# plt.title(f'Grad-CAM Heatmap\nPredicted Class: {predicted_class} ({predicted_probability:.2f})')
# plt.show()


```

```

Prediction Probabilities:
Atelectasis: 0.15
Cardiomegaly: 0.05
Effusion: 0.19
Infiltration: 0.28
Mass: 0.09
Nodule: 0.08
Pneumonia: 0.02
Pneumothorax: 0.08
Consolidation: 0.07
Edema: 0.04
Emphysema: 0.05
Fibrosis: 0.03
Pleural_Thickening: 0.06
Hernia: 0.01

```

Grad-CAM Heatmap
Predicted Class: Infiltration (0.28)



Color Mapping: • The colors in the heatmap correspond to the importance of specific regions of the image: • Red/Hot Colors (e.g., Red, Yellow, Orange): These are areas of high importance or attention for the model. They strongly influenced the model's decision. • Blue/Cool Colors (e.g., Blue, Green): These regions have little to no influence on the model's prediction.

Evaluation

```

# Evaluate the model on the test data
test_loss, test_accuracy = CNN_model.evaluate(test_gen, verbose=1)
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")

# Access ground truth labels
y_true = test_gen.labels # Correct attribute for ground truth labels

# Generate predictions
y_pred_probs = CNN_model.predict(test_gen, verbose=0) # Predicted probabilities
y_pred_classes = np.argmax(y_pred_probs, axis=1) # Predicted class labels

36/36 ————— 19s 516ms/step - accuracy: 0.1196 - loss: 0.1777
Test Loss: 0.1793
Test Accuracy: 0.1176

# Use disease_classes directly for target names
print("\nClassification Report:\n")
print(classification_report(y_true.argmax(axis=1), y_pred_classes, target_names=disease_classes))

```

```

Classification Report:

```

	precision	recall	f1-score	support
Atelectasis	0.00	0.00	0.00	726
Cardiomegaly	0.00	0.00	0.00	20
Effusion	0.00	0.00	0.00	86
Infiltration	0.12	1.00	0.21	132
Mass	0.00	0.00	0.00	35

Nodule	0.00	0.00	0.00	36
Pneumonia	0.00	0.00	0.00	6
Pneumothorax	0.00	0.00	0.00	26
Consolidation	0.00	0.00	0.00	19
Edema	0.00	0.00	0.00	12
Emphysema	0.00	0.00	0.00	10
Fibrosis	0.00	0.00	0.00	9
Pleural_Thickening	0.00	0.00	0.00	4
Hernia	0.00	0.00	0.00	1
accuracy			0.12	1122
macro avg	0.01	0.07	0.02	1122
weighted avg	0.01	0.12	0.02	1122

```

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1531: Und
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1531: Und
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1531: Und
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```

from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

```

```

# Ensure y_true and y_pred_probs are arrays
y_true = test_gen.labels # True labels
y_pred_probs = CNN_model.predict(test_gen, verbose=0) # Predicted probabilities

```

```

# Plot ROC curves for each class
plt.figure(figsize=(10, 8))
for i, class_name in enumerate(disease_classes):
    # Compute ROC curve and AUC for each class
    fpr, tpr, _ = roc_curve(y_true[:, i], y_pred_probs[:, i])
    roc_auc = auc(fpr, tpr)

    # Plot the ROC curve
    plt.plot(fpr, tpr, label=f"{class_name} (AUC = {roc_auc:.2f})")

```

```

# Plot the diagonal (random guess line)
plt.plot([0, 1], [0, 1], 'k--', label="Random Guess")

```

```

# Customize the plot
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for Each Class")
plt.legend(loc="lower right")
plt.show()

```